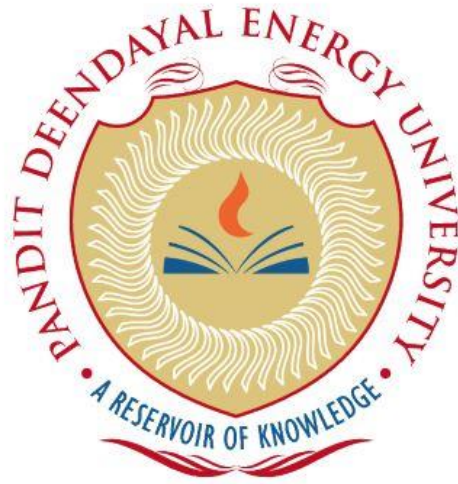


**PANDIT DEENDAYAL ENERGY UNIVERSITY**  
**SCHOOL OF TECHNOLOGY**



**Pattern Recognition Lab**

**20CP412P**

**LAB MANUAL**

**B.Tech. (Computer Science and Engineering)**

**Semester 7**

**Submitted To:**

Dr. Tanmay Bhowmik

**Submitted By:**

HARSH SHAH

21BCP359

G11 Batch

## PRACTICAL 1

<b>Name:</b>	Harsh Shah	<b>Semester:</b>	VII	<b>Division:</b>	6
<b>Roll No.:</b>	21BCP359	<b>Date:</b>	23-07-24	<b>Batch:</b>	G11
<b>Aim:</b>	Calculate the possible eigen values for the given matrix.				

The Eigen Values are:  $3, \frac{7+\sqrt{41}}{2}, \frac{7-\sqrt{41}}{2}$

Q Calculate Possible Eigen Value for matrix A.

Given Matrix :

$$A = \begin{bmatrix} 3 & 0 & 0 \\ 0 & 4 & 5 \\ 0 & 2 & 3 \end{bmatrix}$$

21BCP359
Harsh Shah

Identity Matrix (I) =

$$\begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

calculating  $A - \lambda I =$

$$\begin{bmatrix} 3-\lambda & 0 & 0 \\ 0 & 4-\lambda & 5 \\ 0 & 2 & 3-\lambda \end{bmatrix}$$

To calculate eigenvalues  $|A - \lambda I| = 0$

$$\Rightarrow (3-\lambda) [(4-\lambda)(3-\lambda) - (2)(5)] + 0 + 0 = 0$$

$$\Rightarrow (3-\lambda) [12 - 7\lambda + \lambda^2 - 10] = 0$$

$$= (\lambda - 3) (\lambda^2 - 7\lambda + 2) = 0$$



$$\lambda - 3 = 0$$

$$\boxed{\lambda = 3} \text{ --- (1)}$$

$$\lambda^2 - 7\lambda + 2 = 0$$

$$\lambda = \frac{7 \pm \sqrt{(-7)^2 - 4(1)(2)}}{2}$$

$$\lambda = \frac{7 \pm \sqrt{41}}{2}$$

Thus Eigen values are :  $3, \frac{7+\sqrt{41}}{2}, \frac{7-\sqrt{41}}{2}$

## PRACTICAL 2

<b>Name:</b>	Harsh Shah	<b>Semester:</b>	VII	<b>Division:</b>	6
<b>Roll No.:</b>	21BCP359	<b>Date:</b>	30-07-24	<b>Batch:</b>	G11
<b>Aim:</b>	Extracting Region features and Boundary features from Images				

### Program

```

import requests
from PIL import Image
import numpy as np
import cv2
from io import BytesIO

# List of image URLs
image_urls = [
    "https://images.pexels.com/photos/56866/garden-rose-red-pink-56866.jpeg",
    "https://cdn.pixabay.com/photo/2015/10/09/00/55/lotus-978659_640.jpg",
    "https://s28151.pcdn.co/wp-content/uploads/sites/2/2022/03/Coyote-animal-sentience-research.jpg",
    "https://i.natgeofe.com/k/9acd2bad-fb0e-43a8-935d-ec0aefc60c2f/monarch-butterfly-grass_3x2.jpg",
    "https://image.shutterstock.com/image-photo/green-leaves-philodendron-plant-nature-260nw-2477697533.jpg"
]

# Download images
images = []
for url in image_urls:
    response = requests.get(url)
    img = Image.open(BytesIO(response.content))
    images.append(img)

# Resize images to 256x256 pixels
resized_images = [img.resize((256, 256)) for img in images]

# Convert images to grayscale
gray_images = [cv2.cvtColor(np.array(img), cv2.COLOR_RGB2GRAY) for img in resized_images]

# Extract boundary features using Canny edge detection
boundary_features = [cv2.Canny(img, 100, 200) for img in gray_images]

# Extract region features (using image moments)
region_features = [cv2.moments(img) for img in gray_images]

# Convert region features to a feature vector
feature_vectors = []
for moments in region_features:
    if moments["m00"] != 0:
        cx = int(moments["m10"] / moments["m00"])
        cy = int(moments["m01"] / moments["m00"])

```

```
else:
    cx, cy = 0, 0
    feature_vectors.append([cx, cy])

# Display results
print("Boundary Features (Canny edges):")
for i, bf in enumerate(boundary_features):
    print(f"Image {i+1}:")
    print(bf)

print("\nRegion Features (Centroid coordinates):")
for i, fv in enumerate(feature_vectors):
    print(f"Image {i+1}: Centroid = {fv}")

print("\nFeature Vectors:")
print(feature_vectors)
```

## Output:

```
Region Features (Centroid coordinates):
Image 1: Centroid = [117, 139]
Image 2: Centroid = [120, 124]
Image 3: Centroid = [129, 132]
Image 4: Centroid = [131, 122]
Image 5: Centroid = [130, 131]

Feature Vectors:
[[117, 139], [120, 124], [129, 132], [131, 122], [130, 131]]
```

## PRACTICAL 3

<b>Name:</b>	Harsh Shah	<b>Semester:</b>	VII	<b>Division:</b>	6
<b>Roll No.:</b>	21BCP359	<b>Date:</b>	06-08-24	<b>Batch:</b>	G11
<b>Aim:</b>	Understanding Pre-Processing in Datasets.				

### Question 1

**Dataset:** diabetes.csv

```
import numpy as np
```

```
import pandas as pd
```

```
from sklearn.preprocessing import MinMaxScaler, Binarizer, StandardScaler
```

```
df = pd.read_csv('diabetes.csv')
```

**# Dataset without label/class**

```
df1 = df.drop(['Outcome'], axis=1)
```

**# Scaling**

```
min_max_scaler = MinMaxScaler(feature_range=(0,1))
```

```
scaled_features = min_max_scaler.fit_transform(df1)
```

```
scaled_df = pd.DataFrame(scaled_features, columns=df1.columns)
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age
0	0.352941	0.743719	0.590164	0.353535	0.000000	0.500745	0.234415	0.483333
1	0.058824	0.427136	0.540984	0.292929	0.000000	0.396423	0.116567	0.166667
2	0.470588	0.919598	0.524590	0.000000	0.000000	0.347243	0.253629	0.183333
3	0.058824	0.447236	0.540984	0.232323	0.111111	0.418778	0.038002	0.000000
4	0.000000	0.688442	0.327869	0.353535	0.198582	0.642325	0.943638	0.200000
...	...	...	...	...	...	...	...	...
763	0.588235	0.507538	0.622951	0.484848	0.212766	0.490313	0.039710	0.700000
764	0.117647	0.613065	0.573770	0.272727	0.000000	0.548435	0.111870	0.100000
765	0.294118	0.608040	0.590164	0.232323	0.132388	0.390462	0.071307	0.150000
766	0.058824	0.633166	0.491803	0.000000	0.000000	0.448584	0.115713	0.433333
767	0.058824	0.467337	0.573770	0.313131	0.000000	0.453055	0.101196	0.033333

Figure 1: Scaled df

**# Binarization**

```
binarizer = Binarizer(threshold=0.0)
```

```
binarized_data = binarizer.fit_transform(scaled_df)
```

```
binarized_df = pd.DataFrame(binarized_data, columns=scaled_df.columns)
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age
0	1.0	1.0	1.0	1.0	0.0	1.0	1.0	1.0
1	1.0	1.0	1.0	1.0	0.0	1.0	1.0	1.0
2	1.0	1.0	1.0	0.0	0.0	1.0	1.0	1.0
3	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.0
4	0.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0

Figure 2: Binarized df.head()

**# Standardization**

```
scaler = StandardScaler()
```

```
standardized_data = scaler.fit_transform(binarized_df)
```

```
standardized_df = pd.DataFrame(standardized_data, columns=binarized_df.columns)
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age
0	0.411035	0.080951	0.218515	0.647760	-1.026390	0.120545	0.036108	0.298934
1	0.411035	0.080951	0.218515	0.647760	-1.026390	0.120545	0.036108	0.298934
2	0.411035	0.080951	0.218515	-1.543781	-1.026390	0.120545	0.036108	0.298934
3	0.411035	0.080951	0.218515	0.647760	0.974289	0.120545	0.036108	-3.345217
4	-2.432883	0.080951	0.218515	0.647760	0.974289	0.120545	0.036108	0.298934

Figure 3: Standardized df.head()

---

## Question 2

**Dataset:** spam.csv

```
import re
```

```
import nltk
```

```
import pandas as pd
```

```
from nltk.corpus import stopwords
```

```
nltk.download("stopwords")
```

```
df = pd.read_csv("spam.csv", encoding="latin-1")
```

	v1	v2
0	ham	Go until jurong point, crazy.. Available only ...
1	ham	Ok lar... Joking wif u oni...
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...
3	ham	U dun say so early hor... U c already then say...
4	ham	Nah I don't think he goes to usf, he lives aro...

Figure 4: `df.head()`

### # Remove Punctuation and Stopwords

```
def remove_punctuations(text):
```

```
    return re.sub(r"[^\w\s]", "", text)
```

```
def remove_stopwords(text):
```

```
    stop_words = set(stopwords.words("english"))
```

```
    return " ".join([word for word in text.split() if word.lower() not in stop_words])
```

```
df["v2"] = df["v2"].apply(remove_punctuations)
```

```
df["v2"] = df["v2"].apply(remove_stopwords)
```

	v1	v2
0	ham	Go jurong point crazy Available bugis n great ...
1	ham	Ok lar Joking wif u oni
2	spam	Free entry 2 wkly comp win FA Cup final tkts 2...
3	ham	U dun say early hor U c already say
4	ham	Nah dont think goes usf lives around though

Figure 5: `df.head()`

## PRACTICAL 4

<b>Name:</b>	Harsh Shah	<b>Semester:</b>	VII	<b>Division:</b>	6
<b>Roll No.:</b>	21BCP359	<b>Date:</b>	13-08-24	<b>Batch:</b>	G11
<b>Aim:</b>	Understanding Feature Extraction in Datasets.				

### Question 1

**Dataset:** iris.csv

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
from sklearn.preprocessing import StandardScaler
```

```
from sklearn.decomposition import PCA
```

```
df = pd.read_csv('./Iris.csv')
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa

*# Splitting Features and Target*

```
X = df.drop(['Species'], axis=1)
```

```
y = df['Species']
```

```
X.head()
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
0	1	5.1	3.5	1.4	0.2
1	2	4.9	3.0	1.4	0.2
2	3	4.7	3.2	1.3	0.2
3	4	4.6	3.1	1.5	0.2
4	5	5.0	3.6	1.4	0.2

```
y.head()
```

```
0    Iris-setosa
1    Iris-setosa
2    Iris-setosa
3    Iris-setosa
4    Iris-setosa
Name: Species, dtype: object
```



**# Standard Scaler**

```
scaler = StandardScaler()
```

```
X_standardized = scaler.fit_transform(X)
```

```
X_standardized_df = pd.DataFrame(X_standardized, columns=X.columns)
```

```
X_standardized_df.head()
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
0	-1.720542	-0.900681	1.032057	-1.341272	-1.312977
1	-1.697448	-1.143017	-0.124958	-1.341272	-1.312977
2	-1.674353	-1.385353	0.337848	-1.398138	-1.312977
3	-1.651258	-1.506521	0.106445	-1.284407	-1.312977
4	-1.628164	-1.021849	1.263460	-1.341272	-1.312977

```
X_standardized_df.describe()
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
count	150.000000	1.500000e+02	1.500000e+02	1.500000e+02	1.500000e+02
mean	0.000000	-4.736952e-16	-6.631732e-16	3.315866e-16	-2.842171e-16
std	1.003350	1.003350e+00	1.003350e+00	1.003350e+00	1.003350e+00
min	-1.720542	-1.870024e+00	-2.438987e+00	-1.568735e+00	-1.444450e+00
25%	-0.860271	-9.006812e-01	-5.877635e-01	-1.227541e+00	-1.181504e+00
50%	0.000000	-5.250608e-02	-1.249576e-01	3.362659e-01	1.332259e-01
75%	0.860271	6.745011e-01	5.692513e-01	7.627586e-01	7.905908e-01
max	1.720542	2.492019e+00	3.114684e+00	1.786341e+00	1.710902e+00

**# Principle Component Analysis**

```
pca = PCA(n_components=2)
```

```
principal_components = pca.fit_transform(X_standardized)
```

```
principal_df = pd.DataFrame(principal_components, columns=['PC1', 'PC2'])
```


```
final_df = pd.concat([principal_df, y], axis=1)
```

```
final_df.head()
```

	PC1	PC2	Species
0	-2.816339	0.506051	Iris-setosa
1	-2.645527	-0.651799	Iris-setosa
2	-2.879481	-0.321036	Iris-setosa
3	-2.810934	-0.577363	Iris-setosa
4	-2.879884	0.670468	Iris-setosa

**# Plot**

```
plt.figure(figsize=(8, 6))  
colors = ['red', 'green', 'blue']  
species_names = ['Iris-setosa', 'Iris-versicolor', 'Iris-virginica']  
for species, color in zip(species_names, colors):  
    indices_to_keep = final_df['Species'] == species  
    plt.scatter(final_df.loc[indices_to_keep, 'PC1'],  
                final_df.loc[indices_to_keep, 'PC2'],  
                c=color, s=50, label=species)  
  
# Add labels and title  
plt.xlabel('Principal Component 1')  
plt.ylabel('Principal Component 2')  
plt.title('PCA of Iris Dataset')  
plt.legend()  
plt.grid()
```



## Question 2

**Dataset:** wine.csv

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
df = pd.read_csv('./wine_data.csv')
```

```
df.head()
```

	class_label	alcohol	malic_acid	ash	alcalinity_of_ash	magnesium	total_phenols	flavanoids	nonflavanoid_phenols	proanthocyanins	color_intensity	hue	OD315_of_d
0	1	14.23	1.71	2.43	15.6	127	2.80	3.06	0.28	2.29	5.64	1.04	
1	1	13.20	1.78	2.14	11.2	100	2.65	2.76	0.26	1.28	4.38	1.05	
2	1	13.16	2.36	2.67	18.6	101	2.80	3.24	0.30	2.81	5.68	1.03	
3	1	14.37	1.95	2.50	16.8	113	3.85	3.49	0.24	2.18	7.80	0.86	
4	1	13.24	2.59	2.87	21.0	118	2.80	2.69	0.39	1.82	4.32	1.04	

```
X = df.drop(['class_label'], axis=1)
```

```
y = df['class_label']
```

**# Standardization**

```
scaler = StandardScaler()
```

```
X_standardized = scaler.fit_transform(X)
```

```
X_standardized_df = pd.DataFrame(X_standardized, columns=X.columns)
```


```
plt.figure(figsize=(10, 6))
```

```
X_standardized_df.boxplot()
```

```
plt.xticks(rotation=45, ha='right')
```

```
plt.title('Boxplots for Outlier Analysis (Standardized Data)')
```

```
plt.show()
```



### # Covariance Matrix


```
cov_matrix_standardized = pd.DataFrame(X_standardized, columns=X.columns).cov()
```

```
plt.figure(figsize=(10, 6))
```

```
sns.heatmap(cov_matrix_standardized, annot=True, cmap='coolwarm', fmt=".2f")
```

```
plt.title('Covariance Matrix Heatmap (Standardized Data)')
```

```
plt.show()
```



**# PCA without specifying components**

```
pca = PCA(n_components=None)
```

```
pca.fit(X_standardized)
```

```
plt.figure(figsize=(8, 5))
```

```
plt.scatter(range(1, len(pca.explained_variance_ratio_) + 1), pca.explained_variance_ratio_,  
            label='Variance Ratio', color='blue', alpha=0.6)
```

```
# plt.plot(range(1, len(pca.explained_variance_ratio_) + 1), pca.explained_variance_ratio_)
```

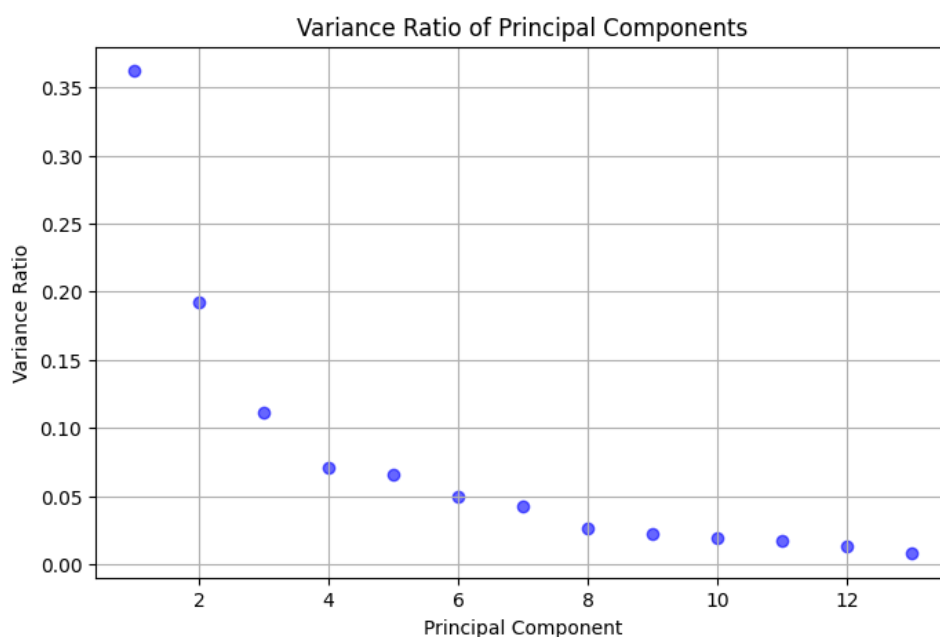
```
plt.xlabel('Principal Component')
```

```
plt.ylabel('Variance Ratio')
```

```
plt.title('Variance Ratio of Principal Components')
```

```
plt.grid()
```

```
plt.show()
```



**# PCA with 2 components**

```
pca_2d = PCA(n_components=2)
```

```
principal_components = pca_2d.fit_transform(X_standardized)
```

```
principal_df = pd.DataFrame(data=principal_components, columns=['PC1', 'PC2'])
```

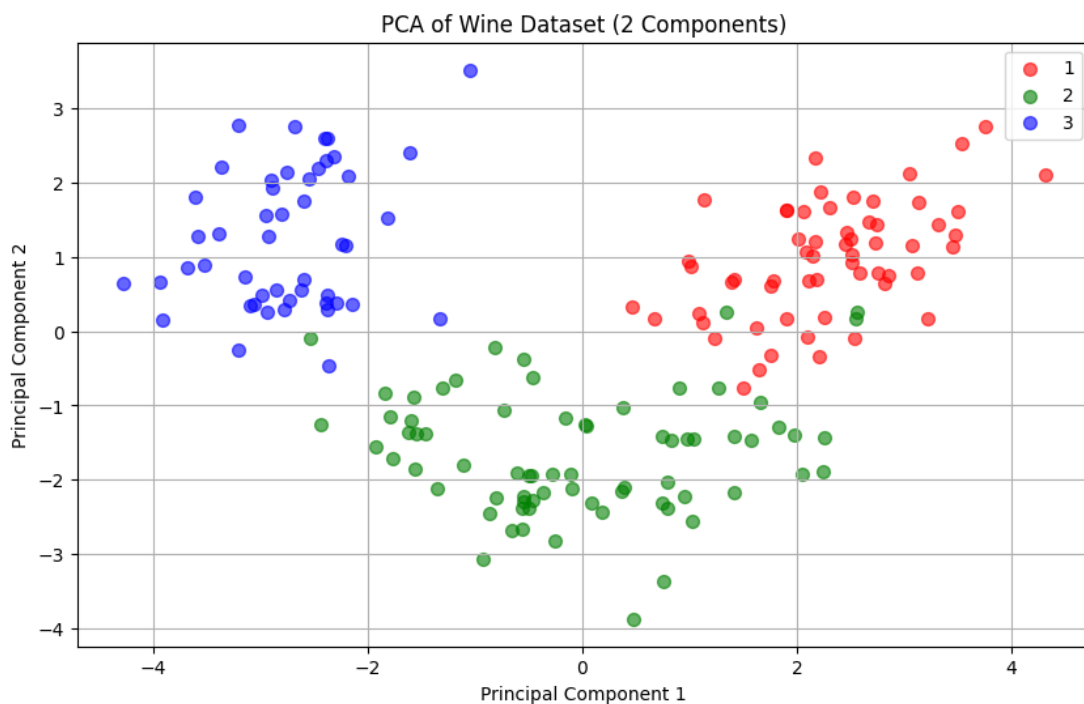
```
final_df = pd.concat([principal_df, y.reset_index(drop=True)], axis=1)
```

```
plt.figure(figsize=(10, 6))

colors = ['red', 'green', 'blue']

for label, color in zip(df['class_label'].unique(), colors):
    indices_to_keep = final_df['class_label'] == label
    plt.scatter(final_df.loc[indices_to_keep, 'PC1'],
               final_df.loc[indices_to_keep, 'PC2'],
               c=color, s=50, label=label, alpha=0.6)

plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.title('PCA of Wine Dataset (2 Components)')
plt.legend()
plt.grid()
plt.show()
```



## PRACTICAL 5

<b>Name:</b>	Harsh Shah	<b>Semester:</b>	VII	<b>Division:</b>	6
<b>Roll No.:</b>	21BCP359	<b>Date:</b>	20-08-24	<b>Batch:</b>	G11
<b>Aim:</b>	Understanding Linear Discriminant projection in Datasets.				

Compute the Linear Discriminant projection for the following two dimensional dataset.

- Samples for class  $\omega_1$ :  $X_1 = (x_1, x_2) = \{(6, 4), (4, 5), (3, 4), (5, 7), (6, 6)\}$
- Sample for class  $\omega_2$ :  $X_2 = (x_1, x_2) = \{(11, 12), (7, 9), (10, 7), (10, 9), (12, 10)\}$

### Linear Discriminant Projection

Linear Discriminant Projection (LDP) refers to the process of projecting data onto a lower-dimensional space in a way that maximizes the separation between different classes. It is a key part of **Linear Discriminant Analysis (LDA)**, a method used in statistics, pattern recognition, and machine learning for dimensionality reduction and classification.

#### Steps:

1. Define the samples for each class
2. Compute the mean vectors
3. Compute the within-class scatter matrix SW for both classes
4. Compute the between-class scatter matrix SB
5. Compute the eigenvalues and eigenvectors of  $SW^{-1} * SB$ 
  - a. First, compute the inverse of SW
  - b. Then, compute the matrix  $SW^{-1} * SB$
  - c. Compute the eigenvalues and eigenvectors
  - d. Find the eigenvector corresponding to the largest eigenvalue

#### Code

```
import numpy as np
```

```
X1 = np.array([[6, 4], [4, 5], [3, 4], [5, 7], [6, 6]]) # Class  $\omega_1$ 
X2 = np.array([[11, 12], [7, 9], [10, 7], [10, 9], [12, 10]]) # Class  $\omega_2$ 
```

```
# Step 1: Compute the mean vectors
```

```
mu1 = np.mean(X1, axis=0)
```

```
mu2 = np.mean(X2, axis=0)
```

```
# Step 2: Compute the within-class scatter matrices
```

```
S_W1 = np.dot((X1 - mu1).T, (X1 - mu1)) / (len(X1) - 1)
```

```
S_W2 = np.dot((X2 - mu2).T, (X2 - mu2)) / (len(X2) - 1)
```

```
S_W = S_W1 + S_W2
```

```
# Step 3: Compute the between-class scatter matrix
```

```
mu_diff = (mu2 - mu1).reshape(2, 1)
```

```
S_B = np.dot(mu_diff, mu_diff.T)
```

```

# Step 4: Compute the projection vector (eigenvector)
eigvals, eigvecs = np.linalg.eig(np.linalg.inv(S_W).dot(S_B))

# Sort eigenvectors by eigenvalues in descending order
eigvecs = eigvecs[:, np.argsort(-eigvals)]
w = eigvecs[:, 0] # Projection vector (corresponding to the largest eigenvalue)

# Output the results
print("Mean vector of class  $\omega_1$ :", mu1)
print("Mean vector of class  $\omega_2$ :", mu2)
print("Within-class scatter matrix S_W:\n", S_W)
print("Between-class scatter matrix S_B:\n", S_B)
print("Projection vector w:", w)

# Project the samples onto the new axis
Y1 = np.dot(X1, w)
Y2 = np.dot(X2, w)

print("Projected samples for class  $\omega_1$ :", Y1)
print("Projected samples for class  $\omega_2$ :", Y2)

```

## Output

```

Mean vector of class  $\omega_1$ : [4.8 5.2]
Mean vector of class  $\omega_2$ : [10.  9.4]
Within-class scatter matrix S_W:
[[5.2 1.8]
 [1.8 5. ]]
Between-class scatter matrix S_B:
[[27.04 21.84]
 [21.84 17.64]]
Projection vector w: [0.82816079 0.5604906 ]

```

```

Projected samples for class  $\omega_1$ : [7.21092712 6.11509614 4.72644475 8.06423812 8.33190831]
Projected samples for class  $\omega_2$ : [15.83565584 10.84154089 12.20504206 13.32602325 15.54283543]

```



## PRACTICAL 6

<b>Name:</b>	Harsh Shah	<b>Semester:</b>	VII	<b>Division:</b>	6
<b>Roll No.:</b>	21BCP359	<b>Date:</b>	03-09-24	<b>Batch:</b>	G11
<b>Aim:</b>	Understanding Principal Component Analysis in Datasets.				

1. Using sklearn library import the digits datasets (i.e. through load\_digits()). It is a 8 x 8 pixel images dataset and they are 64-dimensional. In order to retrieve some of the intuition behind the relationship between these points, make use of PCA to reduce the dimension to a manageable number of dimensions (i.e., 2). After reducing the dimension plot them using scatter plots with cmap.
2. How many number of components will be ideal is an important part of using PCA. Using this data plot the cumulative explained variance ratio as a function of the number of components.
3. Try to reconstruct the data using the largest subset of principal components. The idea behind this is that any components with variance much larger than the effect of the noise should be relatively unaffected by the noise. Add some random noise to the dataset and replot it.

### Code

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import load_digits
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler

# Load the digits dataset
digits = load_digits()
X = digits.data # Feature matrix
y = digits.target # Labels

# Standardize the data (PCA works better on standardized data)
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Apply PCA to reduce the dimensionality to 2 dimensions
pca_2d = PCA(n_components=2)
X_pca_2d = pca_2d.fit_transform(X_scaled)
```

*# Scatter plot of the 2D PCA-reduced data*

```
plt.figure(figsize=(8, 6))
scatter = plt.scatter(X_pca_2d[:, 0], X_pca_2d[:, 1], c=y, cmap='Spectral', edgecolor='k', s=60)
plt.colorbar(scatter)
plt.xlabel('First Principal Component')
plt.ylabel('Second Principal Component')
plt.title('2D PCA of Digits Dataset')
plt.show()
```

*# Plot the cumulative explained variance as a function of the number of components*

```
pca_full = PCA().fit(X_scaled)
plt.figure(figsize=(8, 6))
plt.plot(np.cumsum(pca_full.explained_variance_ratio_), marker='o')
plt.xlabel('Number of Components')
plt.ylabel('Cumulative Explained Variance Ratio')
plt.title('Cumulative Explained Variance Ratio by PCA Components')
plt.grid(True)
plt.show()
```

*# Reconstruct the data using the largest subset of principal components*

```
n_components = 30 # Select the top 30 components
pca_reconstruct = PCA(n_components=n_components)
X_pca_reduced = pca_reconstruct.fit_transform(X_scaled)
X_reconstructed = pca_reconstruct.inverse_transform(X_pca_reduced)
```

*# Add noise to the dataset*

```
noise_factor = 0.5
X_noisy = X_scaled + noise_factor * np.random.normal(size=X_scaled.shape)
```

*# Apply PCA again to the noisy data*

```
X_pca_noisy = pca_2d.fit_transform(X_noisy)
```

*# Scatter plot of noisy PCA-reduced data*


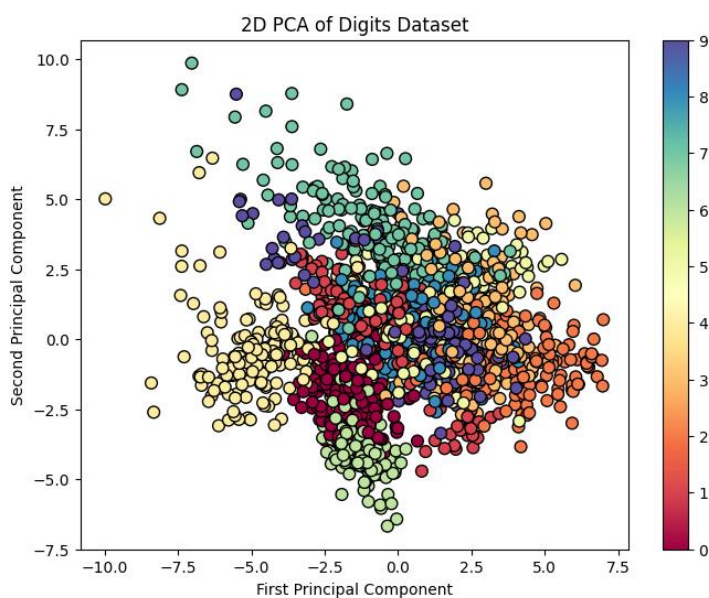
```
plt.figure(figsize=(8, 6))


scatter_noisy = plt.scatter(X_pca_noisy[:, 0], X_pca_noisy[:, 1], c=y, cmap='Spectral', edgecolor='k',
s=60)

plt.colorbar(scatter_noisy)

plt.xlabel('First Principal Component (Noisy)')
plt.ylabel('Second Principal Component (Noisy)')
plt.title('2D PCA of Noisy Digits Dataset')
plt.show()
```

## Output





## PRACTICAL 7

<b>Name:</b>	Harsh Shah	<b>Semester:</b>	VII	<b>Division:</b>	6
<b>Roll No.:</b>	21BCP359	<b>Date:</b>	10-09-24	<b>Batch:</b>	G11
<b>Aim:</b>	Understanding Jaccard Similarity.				

### Jaccard Similarity

Jaccard Similarity is a measure of similarity between two asymmetric binary vectors or we can say a way to find the similarity between two sets. It is a common proximity measurement used to compute the similarity of two items, such as two text documents. The index ranges from 0 to 1. Range closer to 1 means more similarity in two sets of data.

$$J(A, B) = \frac{|A \cap B|}{|A \cup B|}$$

### Code

```
def jaccard_similarity(list1, list2):
```

```
    intersection = 0
    union = 0
```

```
    for a, b in zip(list1, list2):
        if a == 1 or b == 1:
            union += 1
        if a == 1 and b == 1:
            intersection += 1
```

```
    if union == 0:
        return 0
    return intersection / union
```

```
C1 = [0, 1, 0, 0, 0, 1, 0, 0, 1]
```

```
C2 = [0, 0, 1, 0, 0, 0, 0, 0, 1]
```

```
C3 = [1, 1, 0, 0, 0, 1, 0, 0, 0]
```

```
similarity_C1_C2 = jaccard_similarity(C1, C2)
```

```
similarity_C1_C3 = jaccard_similarity(C1, C3)
```

```
similarity_C2_C3 = jaccard_similarity(C2, C3)
```

```
print(f'Similarity - Customer C1 and C2 is {similarity_C1_C2}')
```

```
print(f'Similarity - Customer C1 and C3 is {similarity_C1_C3}')
```

```
print(f'Similarity - Customer C2 and C3 is {similarity_C2_C3}')
```

```
def jaccard_similarity_sets(set1, set2):
```

```
    intersection = len(set(set1).intersection(set2))
```

```
union = len(set(set1).union(set2))  
return intersection / union
```

```
S1 = [0, 2, 5, 7, 9]  
S2 = [0, 1, 2, 4, 5, 6, 8]
```

```
similarity_S1_S2 = jaccard_similarity_sets(S1, S2)
```

```
print(f'Similarity between Set S1 and S2 is {similarity_S1_S2}')
```

## Output

```
Similarity - Customer C1 and C2 is 0.25  
Similarity - Customer C1 and C3 is 0.5  
Similarity - Customer C2 and C3 is 0.0  
Similarity between Set S1 and S2 is 0.3333333333333333
```

## PRACTICAL 8

<b>Name:</b>	Harsh Shah	<b>Semester:</b>	VII	<b>Division:</b>	6
<b>Roll No.:</b>	21BCP359	<b>Date:</b>	10-09-24	<b>Batch:</b>	G11
<b>Aim:</b>	Feature Selection in Dataset.				

### Code

```

import pandas as pd
import matplotlib.pyplot as plt
from statsmodels.tsa.seasonal import seasonal_decompose
from statsmodels.tsa.api import SimpleExpSmoothing
from sklearn.metrics import mean_squared_error

# Step 1: Load the data
data = pd.read_csv("IPG2211A2N.csv", index_col="DATE", parse_dates=True)

# Step 2: Plot the raw data
plt.figure(figsize=(12, 6))
plt.plot(data, label="Industrial Production: Utilities (Electric & Gas)")
plt.title("Industrial Production: Electric & Gas Utilities")
plt.xlabel("Date")
plt.ylabel("Production")
plt.legend()
plt.grid(True)
plt.show()

# Step 3: Trend and Seasonal Variation (Seasonal Decomposition)
decompose_result = seasonal_decompose(
    data, model="multiplicative", period=12
) # Assuming monthly data
decompose_result.plot()
plt.show()

# Step 4: Moving Averages
def plot_moving_average(data, window_sizes):
    plt.figure(figsize=(12, 6))
    plt.plot(data, label="Original", color="blue")

    for window in window_sizes:
        data[f"MA_{window}"] = data["IPG2211A2N"].rolling(window=window).mean()
        plt.plot(data[f"MA_{window}"], label=f"Moving Average {window}-months")

    plt.title("Moving Averages for Industrial Production")
    plt.xlabel("Date")
    plt.ylabel("Production")
    plt.legend()
    plt.grid(True)
    plt.show()

```

```

# Moving averages for 3, 6, and 12 months
plot_moving_average(data.copy(), window_sizes=[3, 6, 12])

# Step 5: Time Series Forecasting
# Using Simple Exponential Smoothing to predict for 2020-2024

# Split data into training and testing
train = data[:"2019"]
test = data["2020":]


# Fit the model on training data
model = SimpleExpSmoothing(train).fit(smoothing_level=0.2, optimized=True)

# Forecast for 2020-2024
forecast = model.forecast(steps=len(test))


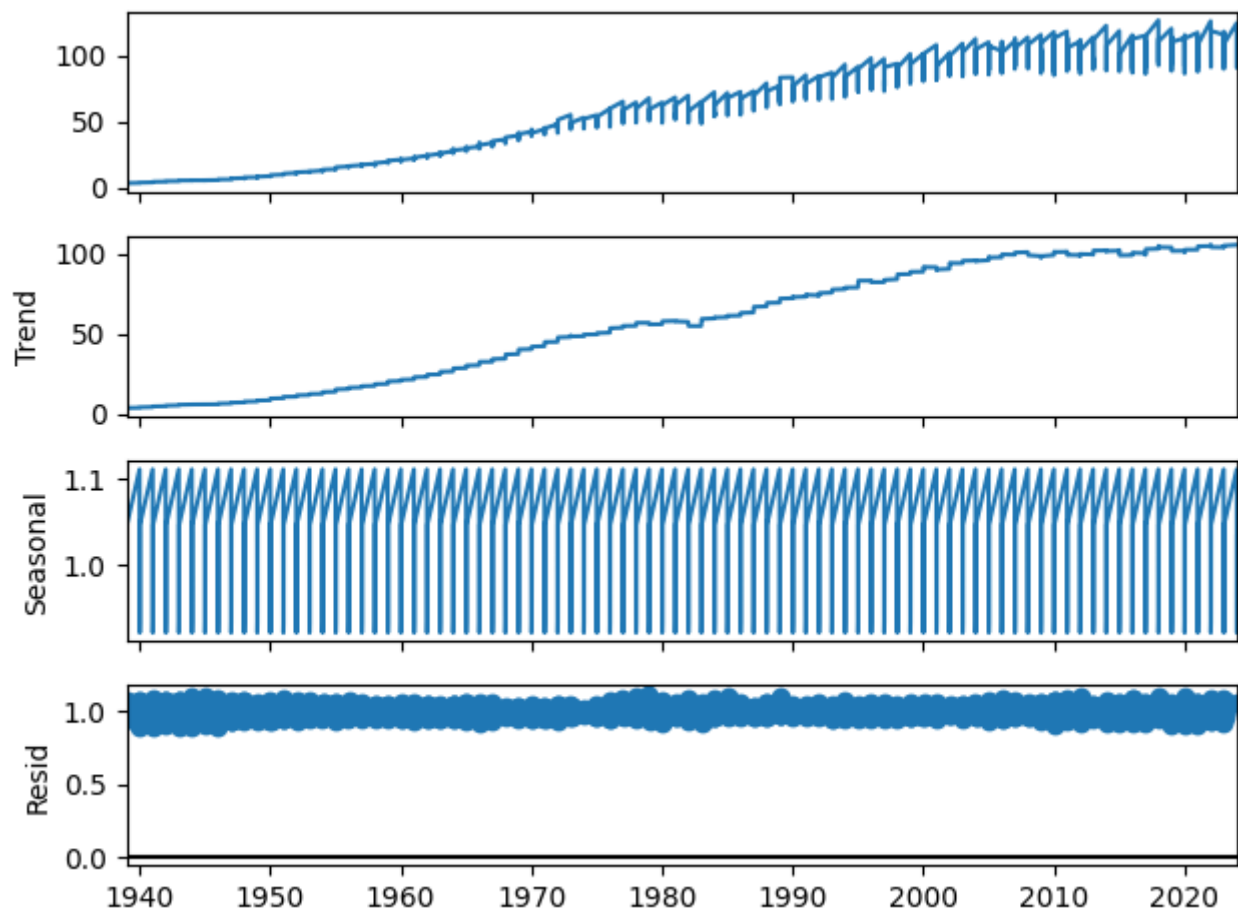
# Plot the forecasted data
plt.figure(figsize=(12, 6))
plt.plot(train, label="Training Data")
plt.plot(test, label="Actual Data (2020-2024)", color="orange")
plt.plot(forecast, label="Forecast (2020-2024)", color="green")
plt.title("Forecasting Industrial Production for Electric & Gas Utilities (2020-2024)")
plt.xlabel("Date")
plt.ylabel("Production")
plt.legend()
plt.grid(True)
plt.show()


# Step 6: Analysis
print(f"Mean Squared Error: {mean_squared_error(test, forecast)}")

```









Mean Squared Error: 107.6594220615686

## PRACTICAL 9

<b>Name:</b>	Harsh Shah	<b>Semester:</b>	VII	<b>Division:</b>	6
<b>Roll No.:</b>	21BCP359	<b>Date:</b>	01-10-24	<b>Batch:</b>	G11

### Code

```
import pandas as pd
import statsmodels.api as sm

df = pd.read_csv('CarPrice_Assignment.csv')

X = df['horsepower']
y = df['citympg']

X = sm.add_constant(X)

# Create the regression model
model_citympg = sm.OLS(y, X).fit()

print(model_citympg.summary())

X_highway = df['horsepower']
y_highway = df['highwaympg']

X_highway = sm.add_constant(X_highway)

# Create the regression model
model_highway = sm.OLS(y_highway, X_highway).fit()

print(model_highway.summary())

import matplotlib.pyplot as plt
import seaborn as sns

# Scatterplot for citympg vs. horsepower
plt.figure(figsize=(8, 6))
sns.scatterplot(x=df['horsepower'], y=df['citympg'], color='blue')
plt.plot(df['horsepower'], model_citympg.predict(sm.add_constant(df['horsepower']))), color='red',
label='Regression Line')
plt.title('Scatterplot of City MPG vs Horsepower')
plt.xlabel('Horsepower')
plt.ylabel('City MPG')
plt.legend()
plt.show()

# Scatterplot for highwaympg vs. horsepower
plt.figure(figsize=(8, 6))
sns.scatterplot(x=df['horsepower'], y=df['highwaympg'], color='green')
```

```
plt.plot(df['horsepower'], model_highway.predict(sm.add_constant(df['horsepower'])), color='red',
label='Regression Line')
plt.title('Scatterplot of Highway MPG vs Horsepower')
plt.xlabel('Horsepower')
plt.ylabel('Highway MPG')
plt.legend()
plt.show()
```

*# Regression Model 1: price as dependent and citympg as independent variable*

```
X_citympg = sm.add_constant(df['citympg'])
model_price_citympg = sm.OLS(df['price'], X_citympg).fit()
```

*# Display model statistics for price vs. citympg*

```
print("Model 1: price vs citympg")
print(model_price_citympg.summary())
```

*# Regression Model 2: price as dependent and highwaympg as independent variable*

```
X_highwaympg = sm.add_constant(df['highwaympg']) # Add constant term
model_price_highwaympg = sm.OLS(df['price'], X_highwaympg).fit()
```

*# Display model statistics for price vs. highwaympg*

```
print("\nModel 2: price vs highwaympg")
print(model_price_highwaympg.summary())
```

*# Model 1: Regression of 'price' on 'enginesize'*

```
X_enginesize = df[['enginesize']]
y_price = df['price']
```

*# Add a constant (intercept)*

```
X_enginesize = sm.add_constant(X_enginesize)
```

*# Create and fit the regression model*

```
model_enginesize = sm.OLS(y_price, X_enginesize).fit()
```

*# Output the model summary*

```
print(model_enginesize.summary())
```

*# Model 2: Regression of 'price' on 'curbweight'*

```
X_curbweight = df[['curbweight']] # Independent variable
y_price = df['price'] # Dependent variable
```

*# Add a constant (intercept)*

```
X_curbweight = sm.add_constant(X_curbweight)
```

*# Create and fit the regression model*

```
model_curbweight = sm.OLS(y_price, X_curbweight).fit()
```

*# Output the model summary*

```
print(model_curbweight.summary())
```

```
# Scatterplot for Engine Size vs. Price
```

```
plt.figure(figsize=(14, 6))
```

```
plt.subplot(1, 2, 1)
```

```
sns.scatterplot(x=df['enginesize'], y=df['price'], color='blue')
```

```
plt.plot(df['enginesize'], model_enginesize.predict(sm.add_constant(df['enginesize'])), color='red',  
label='Regression Line')
```

```
plt.title('Scatterplot of Engine Size vs Price')
```

```
plt.xlabel('Engine Size (L)')
```

```
plt.ylabel('Price ($)')
```

```
plt.legend()
```

```
# Scatterplot for Curb Weight vs. Price
```

```
plt.subplot(1, 2, 2)
```

```
sns.scatterplot(x=df['curbweight'], y=df['price'], color='green')
```

```
plt.plot(df['curbweight'], model_curbweight.predict(sm.add_constant(df['curbweight'])), color='orange',  
label='Regression Line')
```

```
plt.title('Scatterplot of Curb Weight vs Price')
```

```
plt.xlabel('Curb Weight (lbs)')
```

```
plt.ylabel('Price ($)')
```

```
plt.legend()
```

```
# Show the plots
```

```
plt.tight_layout()
```

```
plt.show()
```

```
import pandas as pd
```

```
import statsmodels.api as sm
```

```
from statsmodels.stats.outliers_influence import variance_inflation_factor
```

```
df = pd.read_csv('CarPrice_Assignment.csv')
```

```
# Select numeric variables except 'citympg' and 'highwaympg'
```

```
numeric_df = df.drop(['price', 'citympg', 'highwaympg'], axis=1).select_dtypes(include=[float, int])
```

```
independent_vars = sm.add_constant(numeric_df)
```

```
# Variance Inflation Factor (VIF) calculation
```





```
vif_data = pd.DataFrame()
```

```
vif_data['Feature'] = independent_vars.columns
```

```
vif_data['VIF'] = [variance_inflation_factor(independent_vars.values, i) for i in  
range(independent_vars.shape[1])]
```

```
print(vif_data)
```

# Output



## PRACTICAL 10

<b>Name:</b>	Harsh Shah	<b>Semester:</b>	VII	<b>Division:</b>	6
<b>Roll No.:</b>	21BCP359	<b>Date:</b>	08-10-24	<b>Batch:</b>	G11

### Code

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from statsmodels.tsa.holtwinters import ExponentialSmoothing
from statsmodels.tsa.ar_model import AutoReg
from sklearn.linear_model import LinearRegression

data = pd.read_csv('CarPrice_Assignment.csv')

# Assume data is sequential (e.g., monthly observations)
data['time_index'] = np.arange(len(data))

# Use the 'price' column as the time series target
data['price'] = pd.to_numeric(data['price'], errors='coerce')
data.dropna(subset=['price'], inplace=True)

# Exponential Smoothing Model
exp_model = ExponentialSmoothing(data['price'], seasonal=None, trend=None,
damped_trend=False).fit(smoothing_level=0.5)

# Predict future values
exp_forecast = exp_model.forecast(steps=12)

# Plotting
plt.figure(figsize=(10, 6))
plt.plot(data['time_index'], data['price'], label='Original')
plt.plot(data['time_index'], exp_model.fittedvalues, label='Exponential Smoothing')
plt.plot(range(len(data), len(data) + 12), exp_forecast, label='Forecast', linestyle='--')
plt.legend()
plt.title('Exponential Smoothing Forecast')
plt.show()

# Linear Trend Model
# Fit a linear regression model
linear_model = LinearRegression()
linear_model.fit(data[['time_index']], data['price'])

# Predict values using the model
data['linear_trend'] = linear_model.predict(data[['time_index']])

# Plotting

```

```

plt.figure(figsize=(10, 6))
plt.plot(data['time_index'], data['price'], label='Original')
plt.plot(data['time_index'], data['linear_trend'], label='Linear Trend')
plt.legend()
plt.title('Linear Trend Fit')
plt.show()

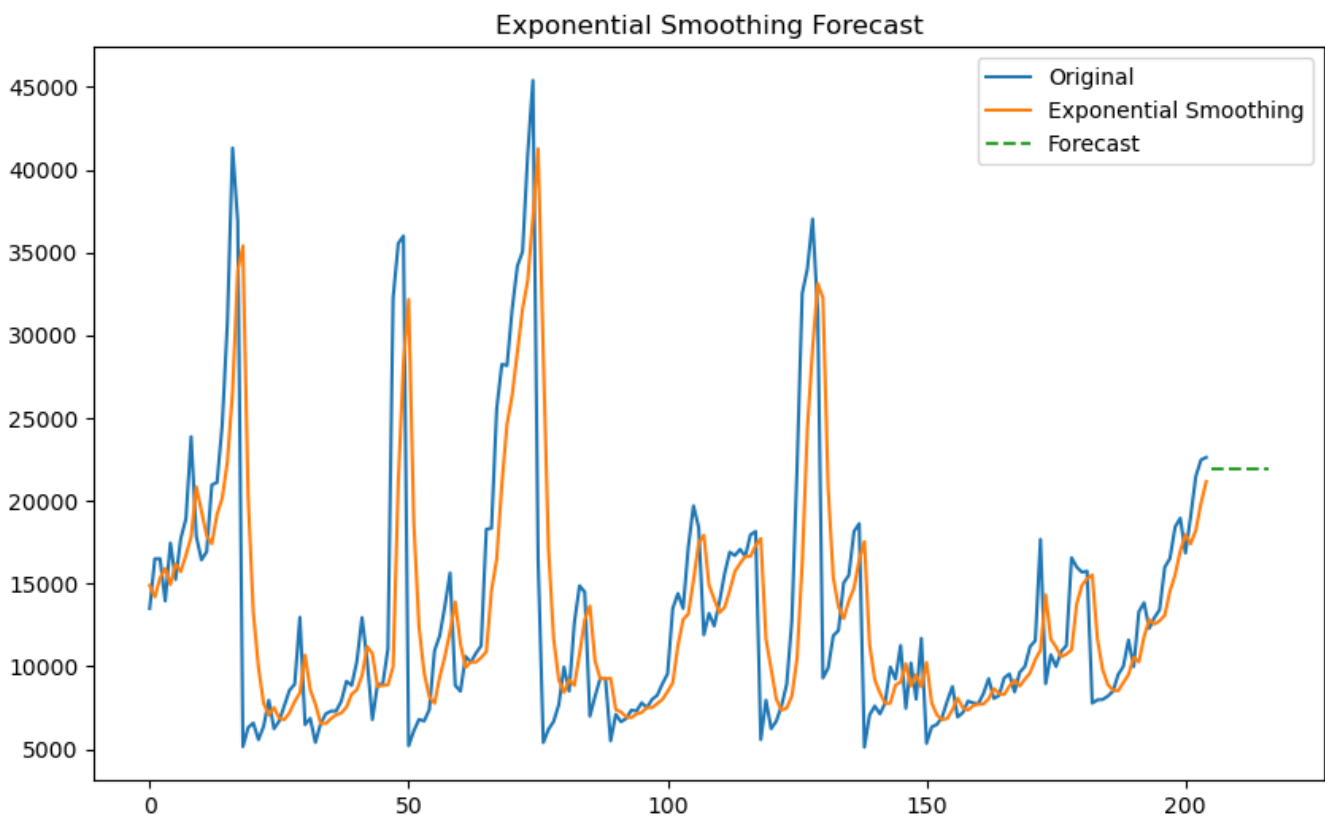
# Autoregressive Model (AR)
# Fit the AR model with a specified lag
ar_model = AutoReg(data['price'], lags=5).fit()

# Predict future values using the AR model
ar_forecast = ar_model.predict(start=len(data), end=len(data) + 11)

# Plotting
plt.figure(figsize=(10, 6))
plt.plot(data['time_index'], data['price'], label='Original')
plt.plot(ar_model.fittedvalues.index, ar_model.fittedvalues, label='AR Fitted Values')
plt.plot(range(len(data), len(data) + 12), ar_forecast, label='AR Forecast', linestyle='--')
plt.legend()
plt.title('Autoregressive Model Forecast')
plt.show()


```

## Output





Linear Trend Fit



Autoregressive Model Forecast

